



Step 3b: Neuron-Level Analysis

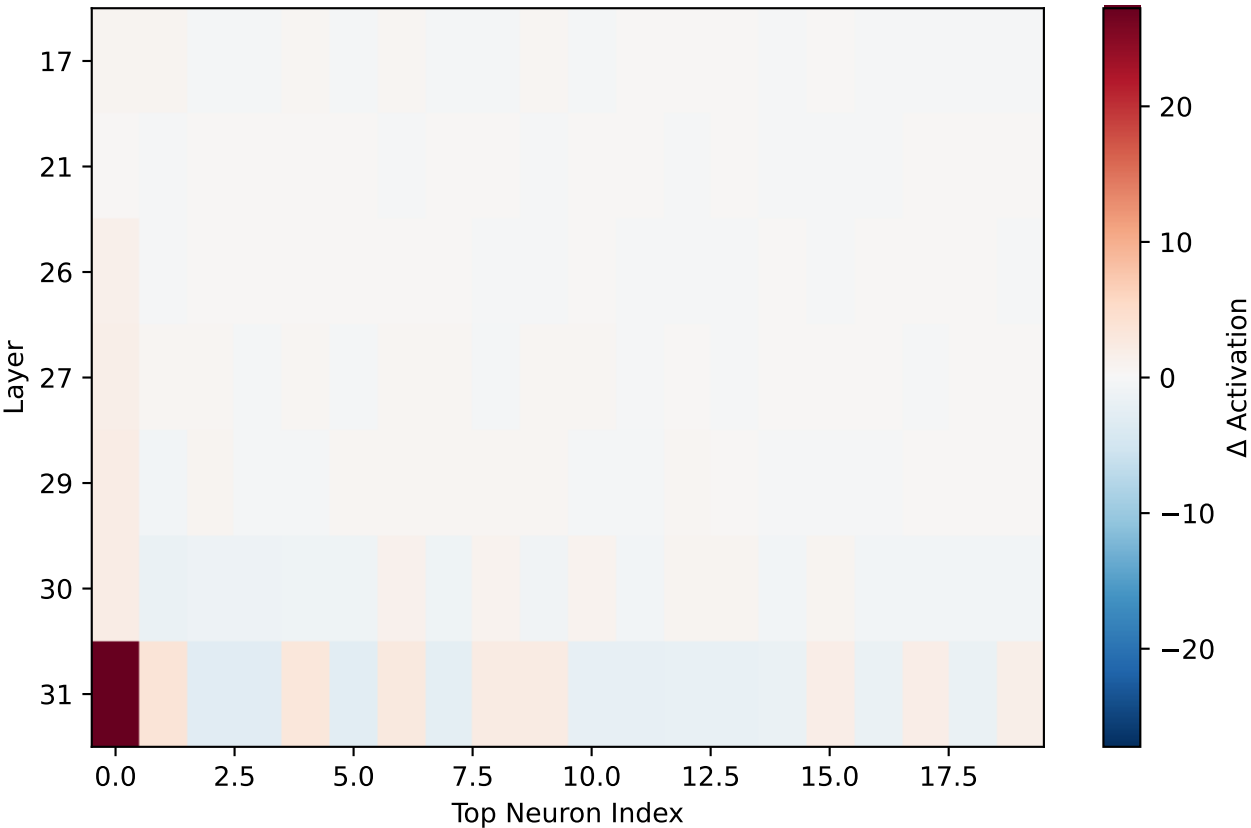
Neuron Activations → Correctness
(Logistic Regression)



Layer 29: Neuron Activation Differences



Top Neurons by Activation Difference



NEURON-LEVEL ANALYSIS SUMMARY

Samples analyzed: 1000
Overall accuracy: 31.3%

Best layer for correctness prediction: Layer 29
- Train accuracy: 0.911
- Test accuracy: 0.863

Key findings:
- Neurons CAN predict correctness (test acc > 50%)
- Some neurons are specifically important for correct vs incorrect predictions
- L1 regularization reveals sparse set of "task neurons"

Top predictive neurons in Layer 29:
- Higher for correct: [1512, 3556, 3554]
- Lower for correct: [2281, 588, 1839]